

Code:

sigmoid.m

```
function g = sigmoid(z)
%SIGMOID Compute sigmoid function
% J = SIGMOID(z) computes the sigmoid of z.

% You need to return the following variables correctly
g = zeros(size(z));

% ===== YOUR CODE HERE =====
% Instructions: Compute the sigmoid of each value of z. (z can be a matrix,
%              vector or scalar).

% add by Chang Liu, for calculating the sigmoid function of each element
g = 1./(1+exp(-z));

% =====

end
```

costFunction.m

```
function [J, grad] = costFunction(theta, X, y)
%COSTFUNCTION Compute cost and gradient for logistic regression
% J = COSTFUNCTION(theta, X, y) computes the cost of using theta as the
% parameter for logistic regression and the gradient of the cost
% w.r.t. to the parameters.

% Initialize some useful values
m = length(y); % number of training examples

% You need to return the following variables correctly
J = 0;
grad = zeros(size(theta));

% ===== YOUR CODE HERE =====
% Instructions: Compute the cost of a particular choice of theta.
%              You should set J to the cost.
%              Compute the partial derivatives and set grad to the partial
%              derivatives of the cost w.r.t. each parameter in theta
%
% Note: grad should have the same dimensions as theta
%

% Chang Liu, for calculating the cost function
temp1 = -1 * (y .* log(sigmoid(X * theta)));
temp2 = (1 - y) .* log(1 - sigmoid(X * theta));
```

```
J = sum(temp1 - temp2) / m;

% for calculating the gradient
% Note: multiple xij is calculate the corresponding value, so we
% use the X' to get the correct value
grad = (X' * (sigmoid(X * theta) - y)) * (1/m);

% =====

end
```

predict.m

```
function p = predict(theta, X)
%PREDICT Predict whether the label is 0 or 1 using learned logistic
%regression parameters theta
% p = PREDICT(theta, X) computes the predictions for X using a
% threshold at 0.5 (i.e., if sigmoid(theta'*x) >= 0.5, predict 1)

m = size(X, 1); % Number of training examples

% You need to return the following variables correctly
p = zeros(m, 1);

% ===== YOUR CODE HERE =====
% Instructions: Complete the following code to make predictions using
% your learned logistic regression parameters.
% You should set p to a vector of 0's and 1's
%

result = sigmoid(X * theta);
for i=1:size(result)
    if result(i) >= 0.5
        p(i) = 1;
    else
        p(i) = 0;
    end
end

% =====

end
```

costFunctionReg.m

```
function [J, grad] = costFunctionReg(theta, X, y, lambda)
%COSTFUNCTIONREG Compute cost and gradient for logistic regression with regularization
% J = COSTFUNCTIONREG(theta, X, y, lambda) computes the cost of using
% theta as the parameter for regularized logistic regression and the
% gradient of the cost w.r.t. to the parameters.
```

```

% Initialize some useful values
m = length(y); % number of training examples

% You need to return the following variables correctly
J = 0;
grad = zeros(size(theta));

% ===== YOUR CODE HERE =====
% Instructions: Compute the cost of a particular choice of theta.
%      You should set J to the cost.
%      Compute the partial derivatives and set grad to the partial
%      derivatives of the cost w.r.t. each parameter in theta

% calculate similar the value of J, theta as before

temp1 = -1 * (y .* log(sigmoid(X * theta)));
temp2 = (1 - y) .* log(1 - sigmoid(X * theta));

% here we add the correction of the theta
thetaT = theta;
thetaT(1) = 0;
correction = sum(thetaT.^2) * (lambda / (2 * m));

% new cost function and gradients by adding correction as equation shows
J = sum(temp1 - temp2) / m + correction;

grad = (X' * (sigmoid(X * theta) - y)) * (1/m) + thetaT * (lambda / m);

% =====

end

```

oneVsAll.m

```

function [all_theta] = oneVsAll(X, y, num_labels, lambda)
%ONEVSALL trains multiple logistic regression classifiers and returns all
%the classifiers in a matrix all_theta, where the i-th row of all_theta
%corresponds to the classifier for label i
% [all_theta] = ONEVSALL(X, y, num_labels, lambda) trains num_labels
% logistic regression classifiers and returns each of these classifiers
% in a matrix all_theta, where the i-th row of all_theta corresponds
% to the classifier for label i

% Some useful variables
m = size(X, 1);
n = size(X, 2);

% You need to return the following variables correctly
all_theta = zeros(num_labels, n + 1);

```

```

% Add ones to the X data matrix
X = [ones(m, 1) X];

% ===== YOUR CODE HERE =====
% Instructions: You should complete the following code to train num_labels
%      logistic regression classifiers with regularization
%      parameter lambda.
%
% Hint: theta(:) will return a column vector.
%
% Hint: You can use y == c to obtain a vector of 1's and 0's that tell use
%      whether the ground truth is true/false for this class.
%
% Note: For this assignment, we recommend using fmincg to optimize the cost
%      function. It is okay to use a for-loop (for c = 1:num_labels) to
%      loop over the different classes.
%
%      fmincg works similarly to fminunc, but is more efficient when we
%      are dealing with large number of parameters.
%
% Example Code for fmincg:
%
%      % Set Initial theta
%      initial_theta = zeros(n + 1, 1);
%
%      % Set options for fminunc
%      options = optimset('GradObj', 'on', 'MaxIter', 50);
%
%      % Run fmincg to obtain the optimal theta
%      % This function will return theta and the cost
%      [theta] = ...
%          fmincg (@(t)(lrCostFunction(t, X, (y == c), lambda)), ...
%              initial_theta, options);
%
for k = 1:num_labels
    init_theta = zeros(n + 1, 1);
    options = optimset('GradObj', 'on', 'MaxIter', 50);
    [theta] = fmincg (@(t)(lrCostFunction(t, X, (y == k), lambda)), init_theta, options);
    all_theta(k,:) = theta;
end;

% =====

end

lrCostFunction.m

function [J, grad] = lrCostFunction(theta, X, y, lambda)
%LRCOSTFUNCTION Compute cost and gradient for logistic regression with

```

```

%regularization
% J = LRCOSTFUNCTION(theta, X, y, lambda) computes the cost of using
% theta as the parameter for regularized logistic regression and the
% gradient of the cost w.r.t. to the parameters.

% Initialize some useful values
m = length(y); % number of training examples

% You need to return the following variables correctly
J = 0;
grad = zeros(size(theta));

% ===== YOUR CODE HERE =====
% Instructions: Compute the cost of a particular choice of theta.
%     You should set J to the cost.
%     Compute the partial derivatives and set grad to the partial
%     derivatives of the cost w.r.t. each parameter in theta
%
% Hint: The computation of the cost function and gradients can be
%     efficiently vectorized. For example, consider the computation
%
%     sigmoid(X * theta)
%
%     Each row of the resulting matrix will contain the value of the
%     prediction for that example. You can make use of this to vectorize
%     the cost function and gradient computations.
%
% Hint: When computing the gradient of the regularized cost function,
%     there're many possible vectorized solutions, but one solution
%     looks like:
%     grad = (unregularized gradient for logistic regression)
%     temp = theta;
%     temp(1) = 0; % because we don't add anything for j = 0
%     grad = grad + YOUR_CODE_HERE (using the temp variable)
%

sig = sigmoid(X * theta);
cost = -y .* log(sig) - (1 - y) .* log(1 - sig);
thetaNoZero = [ 0 ; theta(2:length(theta)) ];
J = (1 / m) * sum(cost) + (lambda / (2 * m)) * sum(thetaNoZero.^ 2);
grad = (1 / m) .* (X' * (sig - y)) + (lambda / m) * thetaNoZero;

% =====

grad = grad(:);

end

predictOneVsAll.m

function p = predictOneVsAll(all_theta, X)
%PREDICT Predict the label for a trained one-vs-all classifier. The labels

```

```
%are in the range 1..K, where K = size(all_theta, 1).
% p = PREDICTONEVSALL(all_theta, X) will return a vector of predictions
% for each example in the matrix X. Note that X contains the examples in
% rows. all_theta is a matrix where the i-th row is a trained logistic
% regression theta vector for the i-th class. You should set p to a vector
% of values from 1..K (e.g., p = [1; 3; 1; 2] predicts classes 1, 3, 1, 2
% for 4 examples)

m = size(X, 1);
num_labels = size(all_theta, 1);

% You need to return the following variables correctly
p = zeros(size(X, 1), 1);


% Add ones to the X data matrix
X = [ones(m, 1) X];

% ===== YOUR CODE HERE =====
% Instructions: Complete the following code to make predictions using
% your learned logistic regression parameters (one-vs-all).
% You should set p to a vector of predictions (from 1 to
% num_labels).
%
% Hint: This code can be done all vectorized using the max function.
% In particular, the max function can also return the index of the
% max element, for more information see 'help max'. If your examples
% are in rows, then, you can use max(A, [], 2) to obtain the max
% for each row.
%

[maxVal, index] = max(X * all_theta', [], 2);
p = index;

% =====
end

lda.m
```



Program paused. Press enter to continue.

Sigmoid function test:

0.000000

0.500000

1.000000

Program paused. Press enter to continue.

Cost at initial theta (zeros): 0.693147

Gradient at initial theta (zeros):

-0.100000

-12.009217

-11.262842

Program paused. Press enter to continue.

Cost at theta found by fminunc: 0.203498

theta:

-25.161272

0.206233

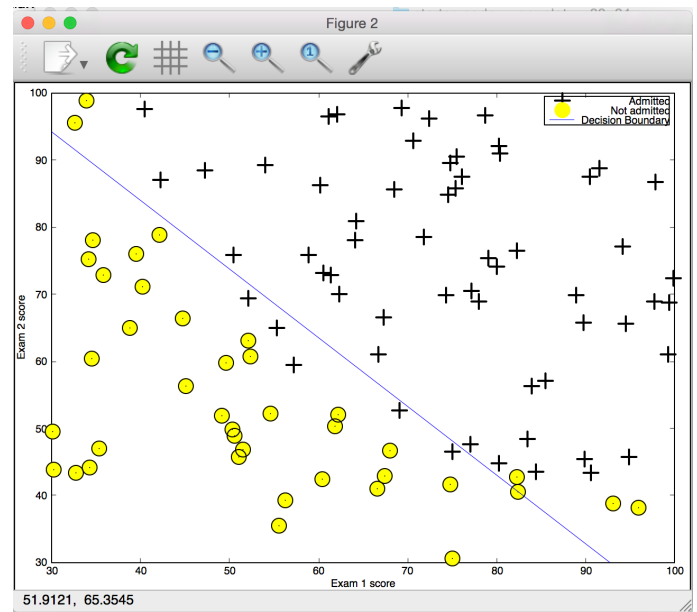
0.201470

Program paused. Press enter to continue.

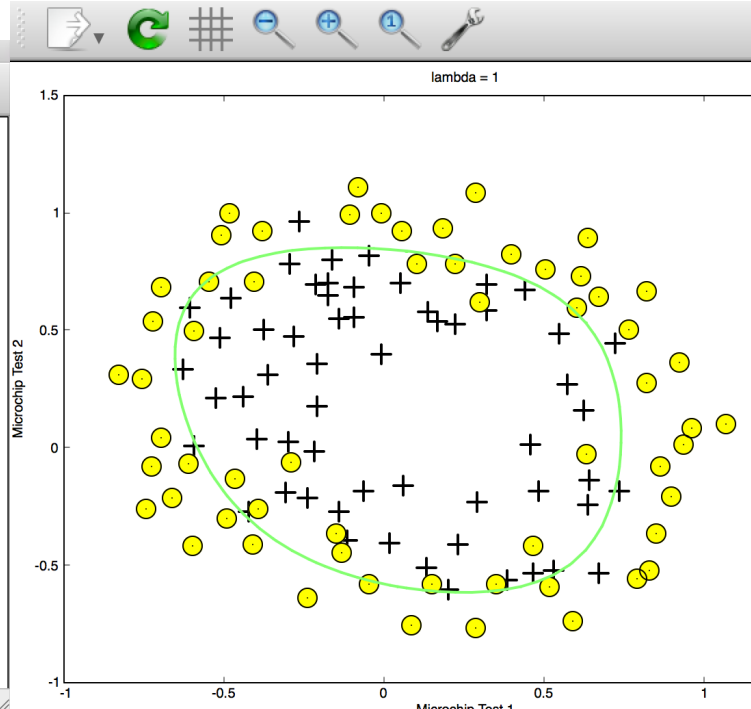

For a student with scores 45 and 85, we predict an admission probability of 0.776

289

Train Accuracy: 89.000000



Screenshot:




gnuplot

Cost at initial theta (zeros): 0.693147

Program paused. Press enter to continue.

Train Accuracy: 83.050847

octave:19>



Loading and Visualizing Data ...
Program paused. Press enter to continue.

Training One-vs-All Logistic Regression

Iteration	50	Cost: 1.401010e-02
Iteration	50	Cost: 5.725249e-02
Iteration	50	Cost: 6.318403e-02
Iteration	50	Cost: 3.589342e-02
Iteration	50	Cost: 6.186099e-02
Iteration	50	Cost: 2.173040e-02
Iteration	50	Cost: 3.557343e-02
Iteration	50	Cost: 8.462181e-02
Iteration	50	Cost: 8.053560e-02
Iteration	50	Cost: 9.611693e-03

Program paused. Press enter to continue.

Training Set Accuracy: 95.100000

octave:20>